

Training at Scale

16M
images

16
datasets

10K+
camera models

48 × A100
GPUs

Indoor: NYUv2, ScanNet, Taskonomy, ...
Outdoor: KITTI, DDAD, Waymo, Lyft, ...
Mixed: DIML, 3D Ken Burns, BlendedMVS, ...

Critical data gap: Only $\sim 20\text{K}$ outdoor images have GT surface normals (mostly from ScanNet indoor). This motivates the depth-normal consistency loss L_{d-n} as a way to learn normals from depth-only outdoor data.