

Key Takeaways

1. **CSTM** — canonical camera transform

2. **Joint optimization** — depth + normals

3. **RPNL** — local depth normalization

Big idea: Transforming the *data space* (CSTM) is simpler and more effective than encoding camera parameters into the network architecture.

Geometric foundation models trained on massive diverse data can rival — and often beat — task-specific models fine-tuned on individual benchmarks.

Questions?