

COMPUTE ALLOCATION REPORT

Microsoft AFMR & Azure for Research

Azure credits for frontier-model and AI-for-science research — OpenAI API access included

J.C. Vaught · April 2026

Program Overview

The Microsoft Research *Accelerating Foundation Models Research* (AFMR) program is the principal channel through which Microsoft distributes Azure compute and Azure OpenAI API quota to academic researchers working on foundation models. Launched in late 2023 and expanded materially through 2024 and 2025, AFMR funds projects in four broad themes. Aligning AI with human goals and values, improving the reliability and safety of foundation models, accelerating scientific discovery in the natural and life sciences, and broadening societal benefit including health, sustainability, and accessibility. Awards are issued as Azure credit grants paired, where relevant, with managed Azure OpenAI Service quota covering GPT-4o, GPT-4.1, the o-series reasoning models, and selected open-weight models hosted on Azure AI Foundry.

The older *Azure for Research* banner historically offered smaller, broader awards distributed through Microsoft Research Outreach. That program has largely been folded into AFMR for foundation-model work and into the Azure sponsorship program for traditional cloud research, so new applicants should default to AFMR. The *Microsoft AI for Good* research lab runs a separate engagement model centered on direct collaboration with Microsoft researchers on humanitarian and sustainability projects, and occasionally co-funds AFMR proposals that align with its mission. AFMR runs on a themed call cadence, with three to four open calls per year, and a permanent public RFP page for general submissions between calls.

Hardware Specs

Azure exposes foundation-model GPUs through the *ND* (high-bandwidth, multi-GPU) and *NC* (single-tenant, smaller) VM families. Pricing varies by region and contract term. The on-demand and one-year reserved figures below are normalized to a per-GPU-hour basis from published Azure list prices for East US 2 and South Central US, where Hopper-class capacity concentrates. Spot pricing typically runs 40–70% below on-demand but is preemptible and not available in every region.

GPUs size	VRAM (GB)	FP16 TFLOPS	BW (TB/s)	On-dem \$/ ter- GPU-hr con- nect	1-yr \$/GPU- hr
ND H200 SXM H200 v5	141	989	4.80	NVLink-4 9.00 900 GB/ s	4.50

GPUs size	VRAM (GB)	FP16 TFLOPS	BW (TB/s)	On-dem ter- GPU-hr con- nect	1-yr \$/GPU- hr
				+ 8 × NDR 400 Gbps	
8 × H100 SXM H100 v5	80	989	3.35	NVLink-4 900 GB/ s + 8 × NDR 400 Gbps	6.98 3.85
8 × A100 (SXM4 80GB)	80	312	2.00	NVLink-3 600 GB/ s + 8 × HDR 200 Gbps	3.67 2.10
4 × A100 RC00 v4	80	312	1.94	PCIe Gen4, no NVLink	3.40 2.00
8 × AMD MI300X v5	192	1307	5.30	In- fin- ity Fab- ric 896 GB/ s + 8 ×	7.20 3.90

GPUs size	VRAM (GB)	FP16 TFLOPS	BW (TB/s)	On-demand \$/GPU-hr	1-yr \$/GPU-hr
				NDR 400 Gbps	

The *MI300X* row is included because AFMR explicitly invites proposals on AMD silicon and Microsoft has been actively building out MI300X capacity in South Central US through 2025–2026. The 192 GB HBM3 per GPU is the largest single-package memory currently available on Azure and is attractive for long-context inference, large mixture-of-experts routing, and activation-heavy training, but the software stack is *ROCm* rather than CUDA and porting effort is non-trivial.

AFMR Allocation Size

AFMR awards are issued as Azure credit balances, sized to the proposal rather than to a fixed tier. Routine thematic-call awards land in the \$10k–\$40k range, with larger flagship calls (notably the science acceleration and safety themes) occasionally reaching \$80k for multi-investigator teams. Azure OpenAI API quota is granted separately as tokens-per-minute and requests-per-minute caps against named deployments, and does not draw down the compute credit balance.

The conversion below maps a dollar credit to *H200-hours* on ND H200 v5, under three pricing scenarios. On-demand uses the \$9/GPU-hr list rate; one-year reserved uses \$4.50/GPU-hr; spot is illustrated at \$2.70/GPU-hr (roughly 70% below on-demand) and is appropriate for checkpointable training runs only.

Credit award	On-demand H200-hr	1-yr reserved H200-hr	Spot H200-hr
\$10,000	1,111	2,222	3,704
\$25,000	2,778	5,556	9,259
\$40,000	4,444	8,889	14,815
\$80,000	8,889	17,778	29,630

A typical \$25k thematic award therefore funds roughly 2,800 H200-hours of on-demand 8-GPU training, or about 5,500 H200-hours under a one-year reservation. For comparison, that is enough to fine-tune a 13B-parameter dense transformer on 50B tokens at industry utilization, or to run several hundred 7B-class supervised fine-tuning sweeps. Spot pricing roughly doubles the budget again for resumable workloads.

Approval Timeline

AFMR runs themed *calls* rather than rolling submission. Each call posts a short description, a deadline (typically four to six weeks out), and a target award range. From submission to award notice is generally four to eight weeks, with the Azure subscription provisioning and OpenAI quota ticket landing one to two

weeks after notification. Researchers should plan for roughly two to three months from initial proposal to first GPU-hour billed.

The proposal itself is lightweight. Two to four pages of research description, a one-paragraph compute and quota budget, short team biographies, and an institutional contact are the standard package. Letters of support are not required. There is no indirect cost recovery because the award is in-kind cloud credit rather than cash.

Eligibility

AFMR is open to academic researchers globally at accredited institutions, which is unusually permissive among foundation-model compute programs. Faculty PIs are the default applicant class, and postdocs may apply with a faculty co-PI. *Graduate students* may serve as PI on certain themed calls when Microsoft explicitly opens the call to students; otherwise they participate as named team members under a faculty PI. All AFMR projects must commit to open publication of results, and use of the Azure OpenAI Service requires accepting Microsoft's data-use and responsible-AI agreements, including content filtering, abuse monitoring, and a prohibition on using outputs to train competing models without separate agreement.

Things to Know

AFMR is currently the most accessible path for an academic researcher to obtain ND H200 v5 capacity together with GPT-4-class API access in a single grant, and the proposal burden is among the lowest in the field. Several practical points materially affect how the award is used. *Awards expire twelve months from activation*, with no rollover, so a burn schedule should be drafted before the subscription is provisioned. *Azure egress is not free*, and a multi-terabyte dataset moved out of region or out of cloud can quietly consume 5–15% of the credit pool; staging datasets in the same region as the compute is essential. *Multi-region ND H200 v5 availability is uneven*, with East US 2 and South Central US the most reliable through 2025–2026 and other regions occasionally queuing for capacity. Microsoft requests, but does not require, that resulting publications acknowledge AFMR by name, and the program team tracks acknowledgments when scoping subsequent calls. The *MI300X* tier is attractive for memory-bound work but locks the codebase to *ROCm*, so proposals should budget engineering time accordingly.

Strategic Fit

For a graduate AI researcher building agentic, evaluation, or foundation-model research, AFMR is best paired with the NAIRR Pilot's Azure pathway. NAIRR delivers a federally reviewed Azure credit through the NAIRR Deep Partnership track, while AFMR delivers a thematically reviewed Azure credit plus Azure OpenAI quota directly from Microsoft Research. Combined, the two programs realistically yield 3,000–5,000 H200-hours per year for a single grad student with a faculty PI, plus substantial GPT-4-class API quota for evaluation harnesses and agent-loop experiments. Where an advisor already holds an active NSF award, *CloudBank* is the right wrapper for cost-managed Azure spend outside the AFMR award. The combination of AFMR (in-kind credit + API), NAIRR (federal Azure pathway), and CloudBank (NSF-

funded paid Azure) covers the full spectrum from prototype to production training without requiring on-prem capacity.

Summary

Tier / program	On-dem H200-hr	1-yr H200-hr	Approval	Best use
AFMR thematic, \$10k	1,111	2,222	4–8 wk	Pilot fine-tunes, eval harnesses
AFMR thematic, \$25k	2,778	5,556	4–8 wk	Mid-scale SFT, RLHF on 7B
AFMR thematic, \$40k	4,444	8,889	4–8 wk	13B-class fine-tune, ablations
AFMR flagship, \$80k	8,889	17,778	6–10 wk	Multi-PI training campaigns
Azure OpenAI quota	n/a	n/a	Bundled	GPT-4o / o-series eval, agents
AI for Good engagement	Variable	Variable	3–6 mo	Humanitarian, sustainability

H200-hour equivalence is computed at dense FP16/BF16 tensor throughput (989 TFLOPS, 141 GB HBM3e, 4.8 TB/s) per H200 SXM GPU. Conversions below use the per-hour FP16 TFLOP ratio; memory-bound workloads scale differently and the bandwidth column is provided so readers can re-derive their own equivalence.

Sources

- AFMR program home and themed calls: microsoft.com/research/collaboration/accelerating-foundation-models-research
- AFMR RFP and submission portal: [AFMR RFP](#) · [AFMR academic program](#)
- Azure ND H200 v5 product page and pricing: [ND H200 v5](#) · [Azure VM pricing](#)
- Azure ND H100 v5 product page: [ND H100 v5](#)
- Azure NDv4 (A100) product page: [NDv4 A100](#)
- Azure ND MI300X v5 announcement: [ND MI300X v5 GA](#)
- Azure OpenAI Service quotas and models: [OpenAI quotas](#) · [OpenAI models](#)
- Microsoft AI for Good Research Lab: [AI for Good](#)
- NAIRR Microsoft Azure contribution: [Microsoft Research NAIRR](#)