

EVOLUTIONARY DISCOVERY OF REINFORCEMENT LEARNING ALGORITHMS VIA LARGE LANGUAGE MODELS

Syggkounas, Loutfi & Persson — GECCO 2026

J.C. Vaught
CSCE 775 — Paper Presentation
University of South Carolina

WHAT RL IS USED FOR

Reasoning LLMs

o1, DeepSeek-R1 — RL makes them think

LLM alignment

RLHF, DPO, GRPO, Constitutional AI

Robotics & autonomy

π_0 , OpenVLA, Atlas, Waymo, AlphaStar

Science & operations

AlphaDev, AlphaTensor, chip design, data-center cooling

This paper targets control RL — Gymnasium / MuJoCo; not LLM-RL.

RL VS. ES AT A GLANCE

Reinforcement Learning

Learn a policy by trial & error, updated with gradients on a scalar reward.

Where it runs today

- AlphaGo, AlphaStar, MuZero
- ChatGPT / Claude / Gemini (RLHF)
- o1, DeepSeek-R1 reasoning
- Humanoid locomotion, manipulation

The dominant paradigm.

Evolutionary Search

Maintain a population, mutate & recombine, keep the fittest. No gradients.

Where it runs today

- Neural architecture search (NAS)
- Hyperparameter tuning (PBT, PB2)
- Reward discovery (Eureka, REvolve)
- Non-differentiable problems (Fun-Search)

The niche tool — used where RL can't go.

THE PROBLEM WITH RL

The same handful of template families keeps getting recycled. LLM alignment has DPO / GRPO / ORPO / SimPO / KTO, but those are all preference-weighted PG. For control RL (this paper's domain), DQN '15, PPO '17, SAC '18, TD3 '18 are still the workhorses.

Sample efficiency

- PPO / SAC: 10^6 – 10^7 env steps on MuJoCo
- Humans: $\sim 10^2$ steps

Reproducibility

- $3\times$ return swings from seeds & HPs
- PPO impls disagree across codebases

Fragility

- Reward hacking, credit-assignment drift
- Deadly triad: bootstrap + FA + off-policy

Human bottleneck

- New algorithm = 1–2 yrs of PhD work
- The update rule has never been automated

THE PROBLEM WITH ES

ES is powerful but boxed in by the designer — flexible and parallel, yet stuck with the search space you wrote down.

Sample cost

- 3–10× more env steps than A3C / PPO
- No gradient, no bootstrap, no TD reuse

Dimensionality

- Stalls past $\sim 10^6$ parameters
- RL scales to billions of params


Search space

- DSL / primitives are hand-designed
- Cannot express what you didn't foresee

Blind mutations

- Random perturbations waste evals
- No semantic guidance on what to change

DEMO --- WHEN RL WINS




Same 20×20 maze, sparse goal reward.

- RL: 212K env steps, ~ 50 -step path
- ES: **$24\times$ more env steps** for the same result

Clean credit assignment is RL's home turf.

DEMO --- WHEN ES WINS



Add two small “trap” rewards near the start.

- RL: **0 / 2,000** episodes reach the big goal
- ES (MAP-Elites): **248 / 1,000** generations do
- 302 total breakthroughs · first at gen 44

Deceptive landscapes are where population search wins.

WHAT THIS PAPER IS

- Each candidate is an executable Python `compute_loss` function
- An LLM proposes mutations and crossovers over that code
- Each variant is scored by actually training agents in Gymnasium
- Extends REvolve — aimed at the **loss function**, not the reward
- Prompt bans actor-critic, value functions, and Bellman updates
- Discovers two new algorithms — CG-FPD and DF-CWP-CP
- Competitive with PPO and SAC across 10 Gymnasium benchmarks

THE CANONICAL RL ALGORITHMS

- DQN (2015) — value-based, discrete action spaces
- TRPO (2015) — trust-region policy gradient
- DDPG (2015) — off-policy continuous control
- PPO (2017) — clipped policy gradient, the workhorse
- SAC (2018) — entropy-regularized off-policy
- TD3 (2018) — twin-critic DDPG

WHAT AUTOML HAS AUTOMATED

- Architecture — NAS, AmoebaNet
- Hyperparameters — PB2, OptFormer
- Reward function — Eureka, REvolve
- Preference loss — DiscoPOP
- The update rule itself — open, until this paper

BACKGROUND

Two lineages converge on one idea

LINEAGE A: DISCOVER AN RL ALGORITHM

EPG '18

neural loss for PG

LPG '20

LSTM update rule

Co-Reyes '21

DAG over typed ops

DiscoRL '25

meta-learned
NN (Nature)

All constrain the search space by hand.

LINEAGE B: LLMS MUTATE CODE

ELM '22

LLMs as mutators

FunSearch '24

cap-set, bin-packing

AlphaEvolve '25


full files,
Strassen beaten

EvoTune '25

proposer fine-tuning

LLMs are **intelligent variation operators** over code.

PARENT WORK: REVOLVE (2024)



What REvolve evolves

- Executable Python **reward functions**
- Island populations, LLM variation
- Tested on driving, humanoid, dexterous hands

The move in this paper.
Evolve the **loss function** instead of the reward.

THE GAP THIS PAPER FILLS

What gets evolved		Representative method
Architecture	✓	NAS, AmoebaNet
Hyperparameters	✓	PB2, OptFormer
Reward function	✓	REvolve, Eureka
Preference loss (align)	✓	DiscoPOP
Update rule (DAG)	✓	Co-Reyes 2021
Update rule (black-box NN)	✓	LPG, DiscoRL
Update rule (executable Python, LLM-evolved)	★	Sygekounas et al. 2026

Everything above is context. The bottom row is **the thesis**.

METHOD

How the framework works

THE BIG PICTURE




Figure 1 from the paper.

Per generation

- LLM proposes candidate loss functions
- Each is trained end-to-end on 5 Gym envs
- Fitness drives selection

Post-evolution

- LLM proposes hyperparameter ranges
- Sweep finds the refined rule $\mathcal{L}^*$

A CANDIDATE IS A CHUNK OF PYTHON

Parent


```
def compute_loss(theta, xi, D):  
    preds = model(D.s, theta)  
    targets = rollout(model, D.s)  
    L = mse(preds, targets)  
    return L
```

After LLM mutation

```
def compute_loss(theta, xi, D):  
    preds = model(D.s, theta)  
    conf = xi.confidence(D.s)  
    targets = rollout(model, D.s)  
    L = (conf * mse(preds, targets)).mean()  
    return L
```



The LLM **edits Python**. That is the entire search space.

THE EVOLUTIONARY LOOP



Island model · 9 individuals per island · 24 new candidates per generation.


WHAT THE LLM IS TOLD NOT TO DO



Constitutional prompt constraints $\Rightarrow$ **forces novel mechanisms.**

Evolution cannot re-invent DQN; it has to find something else.

LLM-GUIDED HPO PASS



Policy arch, optimizer, rollout loop stay frozen. Only internal scalars move.

RESULTS

What evolution discovered

THE TWO ALGORITHMS IT INVENTED

CG-FPD

Confidence-Guided Forward Policy Distillation

Distill short-horizon plans from a learned latent dynamics model. Planning acts as the supervised teacher signal.

No value function. No policy gradient. No Bellman.

DF-CWP-CP


Differentiable Forward Confidence-Weighted Planning with Controllability Prior

Differentiable world-model rollouts + confidence-weighted desirability + controllability prior.

Also avoids every canonical RL mechanism.

Evolution rediscovered **model-based planning** — from scratch.

THE PROPOSER LLM MATTERS. A LOT.



- GPT-5.2 climbs to **0.67**
- Claude 4.5 Opus caps at **0.42**
- Both monotone; GPT-5.2 is sharper on every generation

Same framework, same envs, same seeds — just swap the LLM.


RESULTS MATCH PPO AND SAC

Environment	PPO	SAC	CG-FPD	DF-CWP-CP
CartPole	500.0	—	500.0	500.0
LunarLander	246.6	—	241.2	260.6
MountainCar	−128.1	—	−105.8	−108.6
Acrobot	−63.5	—	−90.6	−78.6
HalfCheetah	1579.1	4988.5	2407.9	2103.8
Swimmer	95.0	87.8	247.5	219.8
Walker2d	3163.2	4595.3	1603.8	1297.6
InvertedPendulum	1000.0	1000.0	1000.0	1000.0

Competitive with hand-designed baselines **across 10 environments**.

Evolved on 5 training envs; remaining 5 are generalization tests.

ABLATION: DIVERSITY WEIGHT



Levenshtein weight α

- $\alpha = 0$: no structural regularization $\Rightarrow 0.50$
- $\alpha = 1$: full regularization $\Rightarrow 0.56$
- $\alpha = 0.5$: **0.65**

Some code-similarity nudge helps;
too much or too little hurts.

DISCUSSION

Critique and open questions

STRENGTHS VS. WEAKNESSES

Strengths

- First **full update-rule** discovery via LLM
- Genuinely novel mechanisms (no value, no PG)
- Clean extension of REvolve's machinery
- LLM-HPO handles fitness fragility cleanly

Weaknesses

- 30 h per candidate on $4\times A100$
- Narrow benchmarks (CartPole / MuJoCo only)
- Is GPT-5.2 discovering or remembering?
- Only 2 evolutionary seeds per LLM
- No head-to-head vs. DiscoRL
- Constitutional constraints never ablated

WHERE THIS PAPER SITS IN THE FIELD

	RL	Alignment
NN-parameterized	DiscoRL (Oh 2025) meta-learned NN	RLHF / DPO (Rafailov 2023) neural reward model
Code	This paper (Sygkounas 2026) evolved Python	DiscoPOP (Lu 2024) LLM-evolved loss

The real contest — **evolved source code** vs. **meta-gradient** for RL.

OPEN QUESTIONS

- Surrogate fitness to break the 30 h/candidate wall.
- Lift the constraints — let evolution decide whether to re-use actor–critic.
- Scale to Atari and robotics — generalization is still untested.
- Curriculum over environments — connect to POET, XLand.
- Self-improving proposer — EvoTune-style DPO on the evolved database.

Takeaway. REvolve showed LLMs can evolve rewards. This paper shows they can evolve the entire algorithm. The open question is whether what they discover is genuinely new — or just a good remix of what the model already saw.

THANK YOU!

Questions & Discussion

jvaught@sc.edu