

RL-TUNED SYNTHETIC MARINE RADAR FOR SIM-TO-REAL TRANSFER




FIGURE 1: Proposal 1 closed-loop method overview.

Method Family	Real-Feedback Adaptive	Env-Type Aware	Sequential Policy Control	Physical Bounds	Task-Reward Optimized
Fixed Domain Randomization (Tobin, Peng)	✗	✗	✗	✗	✗
SimOpt / BayesSim	✓	✗	✗	✗	✗
BayRn / ADR / CDR	✓	✗	✗	✗	✗
Radar Calibration (Digital Twin style)	✓	✗	✗	✓	✗
Proposed Method (Proposal 1)	✓	✓	✓	✓	✓

FIGURE 2: Related-work capability comparison.




FIGURE 3: Published quantitative evidence motivating Proposal 1.

1. INTRODUCTION

Marine radar perception and simulation-driven autonomy have received decades of attention because robust sensing is central to safe maritime robotics and monitoring. However, despite progress in radar signal processing and synthetic-data pipelines, real radar data remain expensive and sparse, and fixed simulator settings still produce a persistent sim-to-real gap that limits downstream detection and mapping performance.

This proposal studies whether a reinforcement learning policy can adaptively tune synthetic radar generator parameters to maximize real-data validation performance across environment types (lake, river, coast). The central hypothesis is that

environment-conditioned simulator control can outperform fixed or random domain randomization and reduce transfer error under unseen routes and weather/clutter conditions.

The approach is grounded in prior work on domain randomization and adaptive sim-to-real transfer [1–13]. However, existing methods either rely on fixed/random randomization, non-RL adaptive tuning, or calibration-centric transfer that is not explicitly conditioned on maritime environment type. Our proposal targets this gap with an environment-aware RL outer loop that directly optimizes real-data downstream task performance, as summarized in Fig. 1 and contrasted with prior work in Fig. 2.

2. APPROACH

We formulate synthetic radar tuning as a bi-level optimization problem. In the inner loop, a downstream model is trained on synthetic radar generated with parameter vector θ . In the outer loop, an RL policy observes the current generator state, environment-type label, and recent real-validation metrics, then outputs updates to θ .

The action space controls key simulator factors including clutter intensity, noise floor, attenuation, speckle, and target reflectivity variation. The reward is defined as real-data task improvement minus realism/stability penalties to discourage unphysical parameter exploitation. We use Soft Actor-Critic (SAC) as the primary outer-loop optimizer due to sample efficiency and continuous-control support, with TD3 as a backup. A non-RL Bayesian optimization tuner is included as a strong baseline.

Relative to existing adaptive randomization methods (SimOpt, BayesSim, BayRn) [3–5], our method makes three concrete changes: (1) sequential policy-based control of simulator updates rather than one-step or purely inference-driven adaptation, (2) explicit conditioning on maritime environment type (‘lake’, ‘river’, ‘coast’), and (3) physically bounded action constraints with realism penalties to prevent simulator exploitation. Relative to radar calibration-oriented sim-to-real pipelines [11], the optimization target is not just simulation fidelity but downstream real-task performance under domain holdout.

3. EVALUATION

Evaluation uses real marine radar logs with explicit environment labels and strict split protocols by route/location to prevent leakage. We compare: (1) fixed synthetic defaults, (2) uniform random domain randomization, (3) Bayesian optimization tuning, (4) RL tuning without environment conditioning, and (5) RL tuning with conditioning.

Primary metrics are downstream real-data task score and relative sim-to-real gap reduction. Secondary metrics include cross-domain generalization (train on two environment types, test on held-out type), unseen-location robustness, seed variance, and performance under low-SNR/high-clutter stress conditions. We will report both best-case performance and sample-efficiency measured by real-validation calls required to reach target accuracy.

To explicitly support the novelty claims, we add claim-linked evaluations: (a) sequential adaptation benefit by comparing RL vs BayRn under matched real-evaluation budgets [5]; (b) environment semantics benefit by comparing conditioned vs unconditioned RL policies; and (c) plausibility benefit by reporting parameter-bound violations and instability with/without realism regularization. These analyses directly test how our method improves over existing randomization and simulator-adaptation approaches [1–4, 6, 7].

Published results already show clear failure modes for non-adaptive baselines (Fig. 3). In real FMCW radar transfer, simulation-only baselines can collapse to near-chance balanced accuracy (about 50% for occupancy and 33% for people-counting), while calibrated adaptation substantially improves performance (about 97% and 72%, respectively) [11]. In synthetic-to-measured SAR transfer, synthetic-only and simple-noise baselines are consistently weaker than stronger adaptation strategies across multiple architectures [10]. These patterns motivate the need for a more structured, environment-aware adaptation controller rather than fixed randomization.

4. CONCLUSION

This project targets a practical and research-relevant problem: improving sim-to-real transfer for marine radar perception through adaptive synthetic data generation. The expected contribution is an RL-controlled simulator tuning framework that is environment-aware, benchmarked against strong non-RL alternatives, and evaluated under realistic transfer stress tests. If full RL tuning is unstable, a contextual-bandit or Bayesian fallback will still provide a complete, meaningful project outcome.

REFERENCES

- [1] Tobin, Josh, Fong, Rachel, Ray, Alex, Schneider, Jonas, Zaremba, Wojciech and Abbeel, Pieter. “Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World.” *arXiv preprint arXiv:1703.06907* (2017).
- [2] Peng, Xue Bin, Andrychowicz, Marcin, Zaremba, Wojciech and Abbeel, Pieter. “Sim-to-Real Transfer of Robotic Control with Dynamics Randomization.” *ICRA* (2018).
- [3] Chebotar, Yevgen, Handa, Ankur, Makoviychuk, Viktor, Macklin, Miles, Issac, Jan, Ratliff, Nathan and Fox, Dieter. “Closing the Sim-to-Real Loop: Adapting Simulation Randomization with Real World Experience.” *arXiv preprint arXiv:1810.05687* (2018).
- [4] Ramos, Fabio, Possas, Rafael and Fox, Dieter. “BayesSim: Adaptive Domain Randomization via Probabilistic Inference for Robotics Simulators.” *arXiv preprint arXiv:1906.01728* (2019).
- [5] Muratore, Fabio, Gienger, Michael and Peters, Jan. “Data-Efficient Domain Randomization with Bayesian Optimization.” *IEEE Robotics and Automation Letters* (2020).
- [6] OpenAI et al. “Solving Rubik’s Cube with a Robot Hand.” *arXiv preprint arXiv:1910.07113* (2019).
- [7] Josifovski, Martin, Cruz, Francisco, Boedecker, Joschka and Wang, Jun. “Continual Domain Randomization.” *arXiv preprint arXiv:2403.12193* (2024).
- [8] Ho, Jonathan et al. “RAP: Learning to Simulate Robust and Accurate Physics for 3D Manipulation.” *arXiv preprint arXiv:2005.07141* (2020).
- [9] Zhao, Wenhao et al. “Sim-to-Real Transfer in Deep Reinforcement Learning for Robotics: A Survey.” *arXiv preprint arXiv:2009.13303* (2020).
- [10] Kim, Wonseok et al. “Soft Segmented Randomization: Enhancing Domain Generalization in SAR-ATR for Synthetic-to-Measured Applications.” *arXiv preprint arXiv:2409.14060* (2024).
- [11] Trinh, Dinh, Deshmukh, Harshal, Herath, Sanjana, Waslander, Steven and Trinh, Hieu. “A Physics-Informed Digital Twin Framework for Calibrated Sim-to-Real FMCW Radar Occupancy Estimation.” *arXiv preprint arXiv:2601.17871* (2026).
- [12] Zust, Luka et al. “LaRS: A Diverse Panoptic Maritime Obstacle Detection Dataset and Benchmark.” *arXiv preprint arXiv:2308.09618* (2023).
- [13] Choi, Jaehyun et al. “PoLaRIS Dataset: A Maritime Object Detection and Tracking Dataset in Pohang Canal.” *arXiv preprint arXiv:2412.06192* (2024).