



How It Actually Works

Reality: state = Python object reference, every state is unique

Episode 1



Episode 2



Q-Table

$(0x3a, \text{restore } 42) = 0.8$	} never accessed again
$(0x3b, \text{restore } 17) = 1.2$	
$(0x3c, \text{remove } 99) = 0.6$	
$(0x3d, \text{restore } 55) = 0.3$	
$(0x7f, \text{restore } 42) = 0.5$	} active now, dead next ep
$(0x80, \text{restore } 88) = 0.9$	
$(0x81, \text{remove } 17) = -0.2$	

After 30 episodes $\times$ 100 steps = 3,000 Q-updates, the table has **3,000 entries keyed on 3,000 unique object refs**. No entry is ever looked up twice. The agent is performing **random search**, not reinforcement learning.